

Package and Function Glossary

- Libraries:
 - cowplot: Add-on to ggplot
 - dplyr: Data wrangling (mutate, select, filter, summarise)
 - readr: Dependency of tidyverse (read.csv, tsv, etc.)
 - tidyverse: Includes readr, dplyr, ggplot among others
 - repr: Create readable text and viewable images of data
 - infer: Used for statistical inference (specify, hypothesize, generate, calculate)
- Main functions:
 - filter()**: Filters rows where conditions specified are true. Separate argument with ,
 - select()**: Returns specific columns. Separate columns with, or add ranges as array [].
 - rep_sample_n()**: Perform repeated sample of size n. can be used to make sample dists.
 - size = size of each sample
 - replace = can be TRUE or FALSE, false by default
 - reps = number of samples
 - pull()**: Pulls a single column from data frame as vector
 - group_by()**: Converts tibble into grouped tibble based on the column value specified.
 - summarise()**: Creates new column with transformed input (ex. mean = mean(column))
 - get_ci()**: Returns a tibble with lower and upper ci
 - level = confidence level
 - specify()**: Specify variable or relationship between variables of interest
 - formula = response ~ explanatory
 - Alternatively, can set response = and explanatory = to variables of choice. Both are needed during hypothesis testing, only response is needed during point estimate
 - success = level of response considered a "success" (used primarily in proportion analysis)
 - hypothesize()**: Declares hypothesis based on variables in specify
 - null = null hypothesis. "independence" is used to determine relationship between response and explanatory. "point" is used to make point estimates
 - mu/med/sigma = true parameter, used with point null hypothesis when response is continuous
 - generate()**: Generates simulated distribution. For CIs, this is a bootstrap distribution. For hypothesis testing, this is a null distribution
 - reps = number of resamples to generate
 - type = method of generating resamples. "bootstrap" used to get bootstrap, "permute" used to get null dist (randomly assigns an in-

- put to a new output in each replicate).
- calculate()**: Returns statistic specified with stat argument
 - stat = type of stat, such as mean, median, sum, sd, prop, diff in means/medians/props
 - order = vector specifying the order in which explanatory variables should be subtracted, ie. c(first, second)

• Sample Plot Workflows:

◦ Histogram using ggplot:

```
plot <- ... %>%
  ggplot(aes(x = ...)) +
  geom_histogram(bins/binwidth =
    ...) +
  ggtitle("...") +
  xlab("...")
```

◦ Boxplot using ggplot:

```
plot <- ... %>%
  ggplot(aes(x = ..., y = ...))
  +
  geom_boxplot() +
  ggtitle("...") +
  ylab("...") +
  geom_hline(yintercept = ...,
    color = "blue")
```

→ obs.stat is the observed stat from the sample/pop

→ direction can be left, right, or two-sided

→ endpoints is a tibble with lower and upper ci

• Sampling Workflows:

◦ Sampling from Population/Bootstrapping from Sample:

```
estimates <- ... %>%
  rep_sample_n(size = ..., reps
    = ..., replace = ...) %>%
  group_by(replicate) %>%
  summarise(...) = ...)
```

→ replace = true for bootstrap, false for sampling

◦ Sampling then Selecting Variable:

```
estimates <- ... %>%
  rep_sample_n(size = ..., reps
    = ..., replace = ...) %>%
  ungroup() %>%
  select(...)
```

◦ Selecting from filtered data:

```
dataset <- ... %>%
  filter(...) = ... %>%
  select(...)
```

◦ Get CI from sample w/ infer package:

```
dataset <- ... %>%
  summarise(lower_ci = quantile(
    stat, ...), upper_ci =
    quantile(stat, ...))
```

Probability and Events

- Sampling w/o replacement yields more precise parameter estimates, sampling w/ replacement guarantees independence.
- Sampling dist. shows all possible values of the sample mean for a given sample size and the likelihood of each value to appear
 - The shapes of the two are not necessarily the same/similar
 - Not all sample means are the same value as the pop mean
 - Sampling dist. is centred at the true population mean (may have difference variance)
 - Increasing sample size doesn't necessarily reduce variance of the sample dist. as variance of random sample is independent of size (it'll probably be closer to the population variance)
 - As the number of sampling reps increases, the distribution becomes more smooth (less missing values/gaps) - you can also use larger binwidths in the histogram
- Point estimates may be influenced by sampling bias (whether intentional or not)
- Standard error: used to quantify variation of point estimates in a sampling dist. (sd of sampling dist.). For bootstrapping, its sd of sds in each replicate, divided by the number of bootstraps.
- Estimator: rv whose dist. is the sampling dist. for a specific sample size and parameter
- Bootstrapping: useful for approximating a sampling dist. when we don't have access to the entire population
 - SD of bootstrap is a decent approximation of the sampling dist., differences between sample dist. and bootstrap are not influenced by sample size
 - It won't always be close to the SD or its corresponding sampling dist. as you might get an unlucky biased sample
 - Taking bootstraps larger than original results in an underestimated SE (artificially narrow), smaller bootstraps yields an overestimate
- Confidence interval: Indicates an X percent change that the true population parameter is between the upper and lower bounds

Simulation-Based Hypothesis Testing

• Sample Plot Workflows:

◦ Bootstrapping w/ Infer:

```
samp_dist_mean <- ... %>%
  specify(response = ...) %>%
  generate(type = "bootstrap",
    reps = 10000) %>%
  calculate(stat = "mean")
```

◦ Point Hypothesis (mean/median/props) using Infer:

```
null_model_infer <- ... %>%
  specify(response = ...) %>%
  hypothesize(null = "point", mu
    = mu_0) %>%
  generate(reps = ...) %>%
  calculate(stat = "mean/median/
    prop")
```

◦ Visualizing Hypothesis Test Results:

```
null_model_vis_infer <- null_model
_infer %>%
  visualize(..., bins/binwidth =
    ...) +
  shade_p_value(obs_stat = obs_
    test_stat, direction = "
    left") +
  xlab("...")
```

◦ obs test stat is the test statistic (in these examples, it's the sample mean)

◦ Sampling from Null Distribution (diff in means/props):

```
null_model <- ... %>%
  specify(formula = explanatory
    ~ response) %>%
  hypothesize(null = "
    independence") %>% #
    independence" is used for
    diffs
  generate(reps = ..., type = "
    permute") %>%
  calculate(stat="diff_in_means/
    props", order = c("mu_1", "
    mu_2"))
```

• Calculate p-value (difference):

```
p_value <- ... %>%
  get_p_value(obs_stat = ...,
    direction = "both")
```

• Sampling dist and Calc Z-score for each Replicate (sample mean):

```
zscore_sample_means <- ... %>%
  rep_sample_n(reps = ..., size =
    ..., replace = FALSE) %>%
  group_by(replicate) %>%
  summarise(sample_mean = mean
    (...)) %>%
  mutate(z = sqrt(n) * (sample_
    mean - mu) / sigma )
```

◦ In this case, mu and sigma are given to us from our initial sample

• Histogram of Z-scores:

```
sampling_dist_sample_mean_z <-
zscore_sample_means %>%
  ggplot() +
```

```
geom_histogram(aes(z, ..density
  ..), color = 'white',
  binwidth = ...) + xlab("...") +
ggtitle("...")
```

• All of the Above + Approximating Pop Statistics:

```
n <- 5
sampling_dist_zscore_s <-
  ... %>%
  rep_sample_n(reps = ..., size =
    n, replace = FALSE) %>%
  group_by(replicate) %>%
  summarise(sample_mean = mean
    (...), sample_sd = sd(...))
  %>%
  mutate(z = sqrt(n) * (sample_
    mean - mu) / sample_sd) %>%
  ggplot() + geom_histogram(aes(z
    , ..density..), color = '
    white', binwidth = ...) +
  xlab("...") + ggtitle("...")
```

• One-Sample t-test + p-value (two-sided):

```
## test stat
test_stat <-
  sqrt(nrow(...)) * (mean(...) -
    mu0) / sd(...)
p_value <- 2 * pt(test_stat, df =
  nrow(...)) - 1, lower.tail =
  FALSE)
```

item Recentering + p-value:

```
samp_dist <- ... %>%
  specify(response = ...) %>%
  generate(type = "bootstrap",
    reps = ...) %>%
  calculate(stat = "mean")
null_model <- samp_dist %>%
  mutate(stat = stat - (mu - mu0))
  %>% #sample_mean - null
p_value <- mean(null_model$stat > mu
  )
```

• Above using t.test():

```
... <- tidy(
  t.test(..., mu = mu0)
)
```

◦ We need a vector/column of data points and mu = mu0

• One-sample z-test (proportion):

```
## This is 1-sided greater
phat <- mean(... == "...")
p0 <- 0.5
test_stat <-
  (phat - p0) / sqrt(p0 * (1-p0)/
    nrow(...))
```

```
p_value <-
  pnorm(test_stat, lower.tail =
    FALSE)
```

• Above using prop.test():

```
\item \textbf{One-sample z-test (
  proportion):}
\begin{lstlisting}[language
  = R]
answer3.2.6 <-
  tidy(prop.test(x = sum(... == "
    ..."),
    n = nrow(...),
    p = 0.5,
    alternative = "...",
    conf.level = ...,
    correct=FALSE))

tidy(
  prop.test(
    x = # the number of successes,
    n = # the number of trials,
    p = # p0 (i.e., the value of p
      under H0),
    alternative = # alternative
      hypothesis: "less", "greater
        ", "two.sided"
    conf.level = # the desired
      confidence level,
    correct = FALSE))
```

• Two-sample t-test (difference of means):

```
## This is two-sided
... <- ... %>%
  filter(...) %>% #if cleaning
    needed
  group_by(...) %>%
  summarise(sample_mean = mean
    (...), sample_var = var(...),
    n = n())
test_stat <- (...$sample_mean[2] -
  ...$sample_mean[1]) /
  sqrt(...$sample_var[2]/...$n[2]
    + ...$sample_var[1]/...$n
      [1])
p_value <- 2 * pt(test_stat, df =
  ..., lower.tail=FALSE)
```

• Above using t.test():

```
... <- tidy(
  t.test(x = ... %>% filter(... ==
    "...") %>% pull(...),
    y = ... %>% filter(... == "...")
      %>% pull(...),
    alternative = "two.sided")
)
```

• two-sample z-test (diff in props):

```
qnts <- ... %>%
```

```
group_by(...) %>%
  count(...) %>%
  mutate(phat = n/sum(n))
n1 <- qnts %>%
  filter(... == "...") %>%
  pull(n) %>%
  sum()
n2 <- qnts %>%
  filter(... == "...") %>%
  pull(n) %>%
  sum()
phat1 <- qnts %>%
  filter(... == "..." & ... == "
    ...") %>%
  pull(phat)
phat2 <- qnts %>%
  filter(... == "..." & ... == "
    ...") %>%
  pull(phat)
phat <- (n1 * phat1 + n2*phat2)/(n1
  + n2)
test_stat <- (phat1 - phat2) / (sqrt
  (phat*(1-phat)*(1/n1 + 1/n2)))
p_value <- 2 * pnorm(test_stat,
  lower.tail = FALSE)
```

• Above using prop.test():

```
tidy(prop.test(x = c(n1*phat1, n2*
  phat2),
  n = c(n1, n2),
  correct = FALSE))
```

• Paired t-test (means of two diff pops):

```
... <- ... %>%
  mutate(d = after.... - before_
    ...)
test_stat <-
  sqrt(nrow(...))*mean(...$d)/sd
    (...$d)
p_value <- pt(test_stat, nrow(...) -
  1, lower.tail = FALSE)
```

• Above using t.test():

```
... <-
  tidy(t.test(x = ...$after...,
    y = ...$before...,
    paired = TRUE,
    alternative = 'greater'))
```

• CI tibbles:

```
...,ci <- tibble(
  lower_ci = qt(0.05, df = nrow
    (...)) - 1) * ..._std_error +
    ..._x_bar,
  upper_ci = qt(0.95, df = nrow
    (...)) - 1) * ..._std_error +
    ..._x_bar,
)
```

```
...,clt_ci <-
  tibble(lower_ci = ..._x_bar +
    qnorm(...) * ..._std_error,
    upper_ci = ..._x_bar -
    qnorm(...) * ..._std_
      error)
```

• Calculating type 1, type 2, and power:

```
n <- ...
..._errors <- tibble(type_I_error =
  0.05,
  type_II_error = 1 - pnorm(qnorm
    (0.05, mu = mu0, sd = .../
      sqrt(n)), mu = mu1, .../sqrt
        (n)),
  power_of_test = pnorm(qnorm
    (0.05, mu0, .../sqrt(n)),
    mu1, .../sqrt(n)))
```

• Finding proportion which reject null at each alpha level:

```
n <- ...
... <- ... %>%
  mutate(test_statistic = sqrt(n)
    * (sample_mean - pop_mean) /
      pop_sd) %>%
  mutate(reject_h0 = abs(test_
    statistic) > qnorm(1-alpha/
      2)) %>%
  group_by(population, alpha) %>%
  summarise(proportion_rejection =
    mean(reject_h0))
```

• ANOVA using aov():

```
anova_results <- aov(formula =
  response ~ explanatory, data =
    ...) %>%
  tidy()
```

```
f_stat <- anova_results$statistic[1]
```

```
anova_pval <- anova_results$p.value
  [1]
```

• Prediction interval (0.9), returns vector:

```
... <- ... %>%
  pull(...) %>%
  quantile(c(0.05, 0.95)) %>%
  unname()
```

• True Coverage Probability, returns vector:

```
... <- ... %>%
  summarise(prop = mean(
    between( ..., ...[1],
      ...[2]))
  ) %>%
  pull(prop)
```

• All of the above + sampling

```
sample_pi <- ... %>%
  rep_sample_n(100, replace =
    FALSE, reps = 1000) %>%
  group_by(replicate) %>%
  summarise(lower = quantile(...,
    0.05),
    upper = quantile(...,
    0.95))
```

```
..._vec <- ...$....
```

```
coverage <- sample_pi %>%
  group_by(replicate) %>%
  summarise(coverage = mean(
    between(..._vec, lower,
      upper)
  )) %>%
  pull(coverage)
```

• Bonferroni Adjustment

```
pval_bonf <- p.adjust(...$p_value,
  method = "bonferroni")
count_bonf <- sum(pval_bonf < 0.05)
```

• BH Adjustment

```
pval_bh <- p.adjust(...$p_value,
  method = "BH")
count_bh <- sum(pval_bh < 0.05)
```

- Pop and sample dists show how some individual data points are distributed, whereas sampling dists show how a statistic (aggregate of points) is distributed for a given sample size n
- Sampling Dist properties:
 - Centered at true population parameter
 - Bell-shaped/normal
 - Becomes narrower and more bell-shaped as n increases
- Standard error is measure of the variability of point estimates in sampling distribution, whereas sd (and variance) are measures of spread
- Test statistic: statistic to use for the test
- Null hypothesis: status quo. If true, the test statistic falls in the sampling distribution under H_0
- Alternative hypothesis: conclusion we wish to make (provided there's evidence to support it)
- p-value: Probability of getting a value at least as extreme as the observed one (left/right bound or two-sided)
 - There is an X chance of observing a test stat at least as extreme as the observed value, provided the null is true
- Significance level: Denoted as α , predetermined value so we reject the null is p-value $\leq \alpha$.
- Law of Large Numbers: LLN states that the sample averages converges to population mean as sample size increases (makes sense as the sample begins to more closely resemble the population)

- Normal (μ, σ) properties:
 - Unimodal and bell shaped
 - Symmetric around mean μ
 - SD σ quantifies spread
 - Z-score is calculated as $Z = \frac{(x-\mu)}{\sigma}$
 - Z-score means a value is Z-score SDs higher/lower than the mean
- CLT: for large sample size n, the sampling dist of the mean or proportion converges to the normal distribution REGARDLESS OF POPULATION
 - $\bar{X} \sim N(\mu, \frac{\sigma}{\sqrt{n}})$ or $\hat{p} \sim N(p, \sqrt{\frac{p(1-p)}{n}})$
- CLT Assumptions:
 - Items in sample are IID
 - Sample size is less than 0.1 of pop
 - Sample must be large enough np and $n(1-p)$ must be greater or equal to 10 for proportions, for means usually 30 is ok
- For means, if CLT is satisfied, we have $\bar{X} \sim N(\mu, \frac{\sigma}{\sqrt{n}})$. You can find this for 0.95 confidence using $qnorm(0.975, \bar{X}, \frac{s}{\sqrt{n}})$
 - $CI(\mu, \alpha) = \bar{x} \pm z_{1-\frac{\alpha}{2}} * \frac{s}{\sqrt{n}}$
- For means, if CLT is satisfied, we have $\hat{p} \sim N(p, \sqrt{\frac{p(1-p)}{n}})$
 - $CI(\hat{p}, \alpha) = \hat{p} \pm z_{1-\frac{\alpha}{2}} * \sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$
- For a difference in means, $\bar{X} - \bar{Y} \sim N(\mu_1 - \mu_2), \sqrt{\frac{\sigma_1^2}{n_1} + \frac{\sigma_2^2}{n_2}}$
 - $CI(\hat{p}, \alpha) = \hat{p} \pm z_{1-\frac{\alpha}{2}} * \sqrt{\frac{s_1^2}{n_1} + \frac{s_2^2}{n_2}}$
- For a difference in proportions, $\hat{p}_1 - \hat{p}_2 \sim N(p_1 - p_2), \sqrt{\frac{p_1(1-p_1)}{n_1} + \frac{p_2(1-p_2)}{n_2}}$
 - $CI(\hat{p}, \alpha) = \hat{p} \pm z_{1-\frac{\alpha}{2}} * \sqrt{\frac{\hat{p}_1(1-\hat{p}_1)}{n_1} + \frac{\hat{p}_2(1-\hat{p}_2)}{n_2}}$
- Standardizing a variable Z converts its distribution to standard normal
 - Apply the following transformation: $Z = \frac{\bar{X} - \mu}{\frac{\sigma}{\sqrt{n}}}$
 - For a proportion: $Z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1-p_0)}{n}}}$
 - Use pnorm on the standardized Z value to get a Z value
- In the event we don't have the population sd, we can use s (sample sd). This adds extra variation.
 - $t = \frac{\bar{x} - \mu_0}{s/\sqrt{n}} \sim t_{n-1}$
 - The tails of the t-dist are heavier to account for the additional variation (for lower degrees of freedom)
 - T-dist is centered about 0 with a dof parameter. In 1-sample, dof = n - 1
 - The higher dof, the closer the t-dist is to N(0, 1)
- Steps for 1 proportion z-test:
 - Formulate null and alternative hypothesis (ex. $H_0 : p = 0.5$ vs $H_1 : p_1 \neq 0.5$)
 - Define test statistic (p)

- Calculate null model $\hat{p} = N(p_0, \sqrt{\frac{p_0(1-p_0)}{n}})$
- Take sample
- Calculate the observed test stat \hat{p}
- Contrast observed test stat with null model by calculating p-value
- Two proportion z-test:
 - Formulate null and alternative hypothesis (ex. $H_0 : p_1 = p_2$ vs $H_1 : p_1 \neq p_2$)
 - Define test statistic $\hat{p}_1 - \hat{p}_2$
 - Calculate null model $\hat{p} = N(0, \sqrt{p * (1-p)(\frac{1}{n_1} + \frac{1}{n_2})})$
 - Take sample x 2 and find \hat{p}
 - Calculate the observed test stat $\hat{p}_1 - \hat{p}_2$
 - Contrast observed test stat with null model by calculating p-value ($2 * pnorm(\hat{p}_1 - \hat{p}_2, 0, \sqrt{p * (1-p)(\frac{1}{n_1} + \frac{1}{n_2})})$)
- Error in hypothesis testing faq:
 - Changing significance level does not affect effect size or overlap between SD and null (error). Location of border separating the null and alternative will change.
 - Increasing sample size reduces chance of type 2 error (power of test increases). The size of the overlap between null and alternative decreases
 - Lowering significance reduces chance of type 1 error, but increases type 2
 - By only reporting p-value, you miss info on effect size and error associated with the statistic
 - Increasing sample size increases power as it narrows the sampling dist, thus reducing overlap error between sampling dist and null
- Remember, var of a prop is $\frac{p(1-p)}{n}$
- ANOVA is used to compared between multiple groups, null is $\mu_1 = \mu_2 = \dots = \mu_n$ (one variable is categorical, the other continuous)
 - ANOVA performs better if variances WITHIN groups are the same
 - If p value is low enough, conclusion is that at least 1 group has a mean that is significantly different from others (presence of difference vs quantifying that difference)
- Prediction interval: Interval when making inference on a new observation (not estimation from previous observations)
 - We take upper and lower quantiles of the population distribution.
- As sample size increases:
 - CIs get narrower as we zero in on true value
 - Prediction intervals remain roughly the same - we are getting better at estimating true quantiles making up the bounds of the PI
- True coverage probability is proportion of the pop that truly falls in the interval. If it's less than the specified probability:
 - More individuals will be outside the interval than

we think

- To capture the TCP, we have to shrink the PI
- The Bonferroni adjustment limits probability of seeing at least one Type 1 error to the specified significance level α to α/k (normally, probability of false positive is α)
 - This mitigates the difference between different types
- BH Adjustment controls the false discovery rate (rate significant features are actually null)